

Here is the code I used in the video, for those who prefer reading instead of or in addition to video.

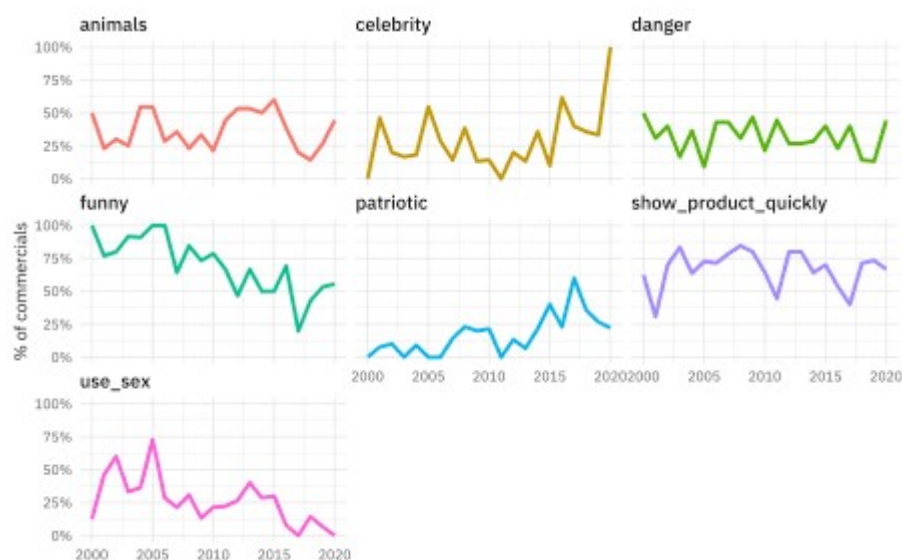
## Explore the data

Our modeling goal is to estimate how the characteristics of [Super Bowl commercials](#) have changed over time. There aren't a lot of observations in this data set, and this is an approach that can be used for robust estimates in such situations. Let's start by reading in the data.

```
library(tidyverse)
youtube <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2021/2021-03-02/youtube.csv")
```

Let's make one exploratory plot to see how the characteristics of the commercials change over time.

```
youtube %>%
  select(year, funny:use_sex) %>%
  pivot_longer(funny:use_sex) %>%
  group_by(year, name) %>%
  summarise(prop = mean(value)) %>%
  ungroup() %>%
  ggplot(aes(year, prop, color = name)) +
  geom_line(size = 1.2, show.legend = FALSE) +
  facet_wrap(vars(name)) +
  scale_y_continuous(labels = scales::percent) +
  labs(x = NULL, y = "% of commercials")
```



## Fit a simple model

Although those relationships don't look perfectly linear, we can use a linear model to estimate if and how much these characteristics are changing with time.

```
simple_mod <- lm(year ~ funny + show_product_quickly +
  patriotic + celebrity + danger + animals + use_sex,
  data = youtube)
```

```

)

summary(simple_mod)

##
## Call:
## lm(formula = year ~ funny + show_product_quickly + patriotic +
##      celebrity + danger + animals + use_sex, data = youtube)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -12.5254  -4.1023   0.1456   3.9662  10.1727
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2011.0838     0.9312  2159.748 < 2e-16 ***
## funnyTRUE       -2.8979     0.8593   -3.372  0.00087 ***
## show_product_quicklyTRUE  0.7706     0.7443    1.035  0.30160
## patrioticTRUE    2.0455     1.0140    2.017  0.04480 *
## celebrityTRUE     2.4416     0.7767    3.144  0.00188 **
## dangerTRUE        0.4814     0.7846    0.614  0.54007
## animalsTRUE       0.1082     0.7330    0.148  0.88274
## use_sexTRUE      -2.4041     0.8175   -2.941  0.00359 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.391 on 239 degrees of freedom
## Multiple R-squared:  0.178, Adjusted R-squared:  0.1539
## F-statistic: 7.393 on 7 and 239 DF, p-value: 4.824e-08

```

We get statistical properties from this linear model, but we can use bootstrap resampling to get an estimate of the variance in our quantities. In tidymodels, the [rsample](#) package has functions to create resamples such as bootstraps.

```

library(rsample)
bootstraps(youtube, times = 1e3)

## # Bootstrap sampling
## # A tibble: 1,000 x 2
##   splits      id
##   <chr>    <chr>
## 1 Bootstrap0001
## 2 Bootstrap0002
## 3 Bootstrap0003
## 4 Bootstrap0004
## 5 Bootstrap0005
## 6 Bootstrap0006
## 7 Bootstrap0007
## 8 Bootstrap0008
## 9 Bootstrap0009
## 10 Bootstrap0010
## # ... with 990 more rows

```

This has allowed you to carry out [flexible bootstrapping or permutation steps](#). However, that's a lot of steps to get to confidence intervals, especially if you are using a really simple model! In a recent release of `rsample`, we added a new function `reg_intervals()` that finds confidence intervals for models like `lm()` and `glm()` (as well as models from the [survival](#) package).

```
set.seed(123)
youtube_intervals <- reg_intervals(year ~ funny + show_product_quickly +
  patriotic + celebrity + danger + animals + use_sex,
  data = youtube,
  type = "percentile",
  keep_reps = TRUE
)
```

```
youtube_intervals
```

```
## # A tibble: 7 x 7
##   term                .lower .estimate .upper .alpha .method
##   .replicates
##
## 1 animalsTRUE        -1.22      0.144  1.51    0.05 percentile
##   [2,001 x 2]
## 2 celebrityTRUE       0.828      2.46   4.06    0.05 percentile
##   [2,001 x 2]
## 3 dangerTRUE         -1.01      0.515  2.09    0.05 percentile
##   [2,001 x 2]
## 4 funnyTRUE          -4.58     -2.91  -1.26    0.05 percentile
##   [2,001 x 2]
## 5 patrioticTRUE      0.112      2.05   3.88    0.05 percentile
##   [2,001 x 2]
## 6 show_product_quicklyT... -0.839      0.740  2.23    0.05 percentile
##   [2,001 x 2]
## 7 use_sexTRUE        -4.04     -2.43  -0.952   0.05 percentile
##   [2,001 x 2]
```


All done!

## Explore bootstrap results

We can use visualization to explore these results. If we had *not* set `keep_reps = TRUE`, we would only have the intervals themselves and could a plot such as this one.

```
youtube_intervals %>%
  mutate(
    term = str_remove(term, "TRUE"),
    term = fct_reorder(term, .estimate)
  ) %>%
  ggplot(aes(.estimate, term)) +
  geom_vline(xintercept = 0, size = 1.5, lty = 2, color = "gray80") +
  geom_errorbarh(aes(xmin = .lower, xmax = .upper),
    size = 1.5, alpha = 0.5, color = "midnightblue"
  ) +
  geom_point(size = 3, color = "midnightblue") +
```

```
labs (
  x = "Increase in year for each commercial characteristic",
  y = NULL
)
```



Since we did keep the individual replicates, we can look at the distributions.

```
youtube_intervals %>%
  mutate(
    term = str_remove(term, "TRUE"),
    term = fct_reorder(term, .estimate)
  ) %>%
  unnest(.replicates) %>%
  ggplot(aes(estimate, fill = term)) +
  geom_vline(xintercept = 0, size = 1.5, lty = 2, color = "gray50") +
  geom_histogram(alpha = 0.8, show.legend = FALSE) +
  facet_wrap(vars(term))
```

